

Assignment 1

Joram Keijser

Monday, November 17, 2014















1, 2 Data

Get the data, divide into test and training set.

```
Data <- read.table("http://www.timvanerven.nl/wp-content/uploads/teaching/statlearn2014/housing-p.data")
Train <- Data[1:350, ]
Test <- Data[351:506, ]
FOLDS = 5
```

There are 13 explanatory variables, from which we try to predict 1 response variable: MEDV, the median value of owner-occupied homes in \$1000's. We plot the each of the 13 explanatory variables against the number of MEDV, the median value of owner-occupied homes in \$1000's. We added the single variable least squares fit to each plot. The independent variables are, on average, influencing MEDV in the way we expect them to. For example: higher criminality leads to a lower price, whereas a higher number of rooms leads to an higher price. Some plots show a nice 'continuous' relationship between prediction and response. See e.g. LSTAT, the percentage lower status of the population. Data on other variables is rather discrete of nature. Examples are RAD, an index of accessibility to radial highways and the binary CHAS. Finally, we notice some variables which do seem to correlate with MEDV in a particular way, but whose values are lumped together around a certain value. See for example CRIM and B. Also included is an histogram showing the distribution of MEDV.

```
par(mfrow = c(5, 3))
for(i in 1:13) {
  x <- Data[, i]
  y <- Data[, 14]
  plot(x, y, xlab = names(Data)[i], ylab = names(Data)[14])
  lm.model <- lm(y ~ x)
  lines(x, lm.model$coefficients[[1]] + lm.model$coefficients[[2]]*x, col = 'red')
}
truehist(Data$MEDV, xlab = 'MEDV', ylab = 'Percentage')
```



3, 4 Regression

Of the 13 explanatory variables, 3 have not-statistically significant regression coefficients in our multivariate linear model. These are **INDUS** (proportion of non-retail business acres per town), **CHAS** (boolean indicating whether the tract bounds on the Charles River) and **AGE** (proportion of owner-occupied units built prior to 1940). The obtained coefficients are not very surprising. The independent variables are, on average, influencing **MEDV** in the way we expect them to. For example: higher criminality leads to a lower price, whereas a higher number of rooms leads to an higher price. In the single variable model based on all data, the variables **INDUS** and **DIS** are correlated negatively and positively, respectively. In the multivariate model build on the training data, these relationships are switched.

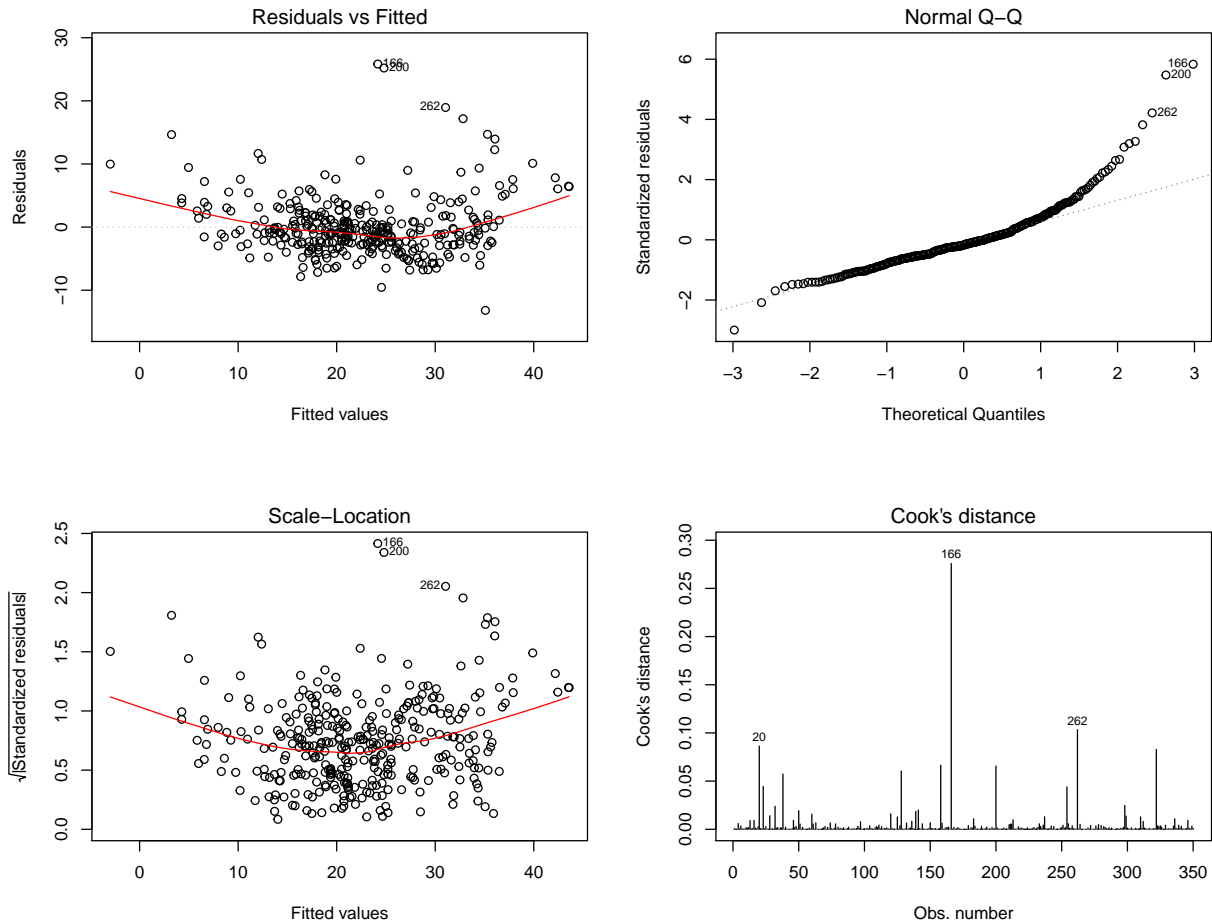
```
lm.model <- lm(MEDV ~ ., data = Train)
summary(lm.model)

##
## Call:
## lm(formula = MEDV ~ ., data = Train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -13.191  -2.659  -0.788   1.732  25.831
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  39.40033     6.17442   6.38 5.8e-10 ***
## CRIM         -0.11352     0.04178  -2.72 0.00692 **
## ZN           0.06372     0.01638   3.89 0.00012 ***
## INDUS        0.02483     0.07374   0.34 0.73647
## CHAS         1.48609     0.99799   1.49 0.13740
## NOX        -17.03198     4.70072  -3.62 0.00034 ***
## RM           3.35805     0.52239   6.43 4.4e-10 ***
## AGE         -0.00537     0.01562  -0.34 0.73101
## DIS        -1.64248     0.23896  -6.87 3.1e-11 ***
## RAD          0.29926     0.07541   3.97 8.8e-05 ***
## TAX        -0.01378     0.00432  -3.19 0.00155 **
## PTRATIO    -0.84220     0.15372  -5.48 8.4e-08 ***
## B           0.00631     0.00318   1.99 0.04770 *
## LSTAT      -0.52157     0.06067  -8.60 3.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.67 on 336 degrees of freedom
## Multiple R-squared:  0.745, Adjusted R-squared:  0.735
## F-statistic: 75.4 on 13 and 336 DF, p-value: <2e-16
```

The plots Residuals vs fitted and Scale-location plot show that the residuals are to some extent, but not entirely randomly distributed when plotted against the fitted values. The Q-Q-plot shows that the residuals are neatly normally distributed except in the very high (and very low) range. A plot of Cook's distance shows that there are 3 clear outliers, but no single observation influences the regression too much.

```
http://stats.stackexchange.com/questions/58141/interpreting-plot-lm
par(mfrow = c(2, 2))
for(i in 1:4) {
```

```
plot.lm(lm.model, which = i)
}
```



5. Cross-validation

The training error equals 7341.966. This corresponds to an average squared difference between predicted and actual MEDV of 20.977. Since the variable is measured in \$1000's, this corresponds to 20977 Dollars. We will see that the cross validation error is higher.

```
performCV <- function(Data, K = FOLDS, fitFUN = lm, ...){
  set.seed(123)
  # Divide data in K folds
  folds <- sample.int(K, size = dim(Data)[1], replace = TRUE)
  # Vector will contain mean error of every fold
  errors <- rep(0, 5)
  aantal <- rep(0, 5)
  for(k in 1:K) {
    # test on fold k, train on other folds:
    aantal[k] <- sum(folds == k)
    test = Data[folds == k, ]
  }
}
```

```

train = Data[folds != k, ]
model <- fitFUN(MEDV ~ ., data = train)
predict <- data.matrix(cbind(1, test[, -dim(test)[[2]]]))%*%data.matrix(model$coef)
errors[k] <- sum((predict - test[, dim(Data)[[2]]])^2)
}
#return(mean(errors)) #averaged over all folds
return(data.frame(errors = errors, samples = aantal))
}
result <- performCV(Train)
total <- sum(result[,1])

```


The mean cross validation error equals 8578.361. Averaged: 24.51, indeed higher than the training error. This does not come as a surprise, since the cross-validation error is supposed to be a more accurate estimate of the model's performance.

6. Full subset search

```

set.seed(123)
subModels <- regsubsets(as.matrix(Train[, -14]), Train[, 14], nbest = 1, nvmax = 20)
LMerrors <- rep(0, dim(summary(subModels)$which)[1])
for( k in 1:length(LMerrors)) {
  select <- as.logical(summary(subModels)$which[k,])
  select[14] <- TRUE
  subset <- Train[, select]
  result <- performCV(subset)
  LMerrors[k] <- sum(result[,1])/sum(result[,2])
}
plot(LMerrors, xlab = 'Number of features', ylab = 'mean error')

```



```
summary(subModels)$which[which(LMerrors == min(LMerrors)), ]
```


```
## (Intercept)      CRIM      ZN      INDUS      CHAS      NOX
##      TRUE      TRUE      TRUE      FALSE      TRUE      TRUE
##      RM      AGE      DIS      RAD      TAX      PTRATIO
##      TRUE      TRUE      TRUE      TRUE      TRUE      TRUE
##      B      LSTAT
##      TRUE      TRUE
```

Minimum of mean error: 23.995, achieved with all 12 features `INDUS`. This error is slightly higher than that achieved when using all 13 features (see figure). This is not a very big surprise, since coefficient of the excluded feature has a very high p-value in the multivariate linear model.


7. Ridge

We use the package `glmnet` which provides higher level functions for performing ridge and lasso regression. The first plot shows that for small values of (log) `lambda`, the coefficients of variables `CHAS`, `NOX` and `RM` (4, 5, and 6, respectively) are very large. Just like we saw in the least squares fit, of course.

```
set.seed(123)
x <- as.matrix(Train[, c(1:13)])
y <- as.matrix(Train[, 14])
fit.ridge <- glmnet(x, y, alpha = 0)
plot(fit.ridge, xvar="lambda", label = TRUE)
```



```
cv.ridge = cv.glmnet(x, y, alpha = 0, nfold = FOLDS)
plot(cv.ridge)
```



```
cv.ridge = cv.glmnet(x, y, alpha = 0, nfolds = FOLDS, lambda = seq(0.001, 2, by = 0.001))
predict <- predict(cv.ridge, newx = x, s = "lambda.min")
minIndex <- which(cv.ridge$lambda == cv.ridge$lambda.min)
Ridge = as.matrix(coef(cv.ridge, s = "lambda.min"))
LM = summary(lm.model)$coef[,1]
```

From the second plot, we notice that the cross-validation error is smallest for values of lambda smaller than 1. By subsequently letting lambda range between 0 and 1, we find that the minimal mean-squared error equals 23.602. The corresponding lambda equals 0.162. This cv-error is less than that of 0.369, achieved by ordinary least squares. The coefficients are indeed smaller than in the LM model. The difference is only -0.071, but the different scales of the prediction variables make this number not directly interpretable.

```
data.frame(Ridge = Ridge, LM = LM)
```


##	X1	LM
## (Intercept)	36.03344	39.40033
## CRIM	-0.10405	-0.11352
## ZN	0.05756	0.06372
## INDUS	-0.00258	0.02483
## CHAS	1.62462	1.48609
## NOX	-15.02044	-17.03198
## RM	3.48789	3.35805
## AGE	-0.00612	-0.00537
## DIS	-1.51390	-1.64248
## RAD	0.23958	0.29926


```
## TAX          -0.01095 -0.01378
## PTRATIO      -0.81602 -0.84220
## B            0.00641  0.00631
## LSTAT        -0.50413 -0.52157
```


8 Lasso

Also using `glmnet` package.

```
set.seed(123)
fit.lasso<- glmnet(x, y, alpha = 1)
plot(fit.lasso, xvar="lambda", label = TRUE)
```



```
cv.lasso = cv.glmnet(x, y, alpha = 1, nfolds = FOLDS, lambda = seq(0.001, 2, by = 0.001))
plot(cv.lasso)
```



```
predict <- predict(cv.lasso, newx = x, s = "lambda.min")
minIndex <- which(cv.lasso$lambda == cv.lasso$lambda.min)
Lasso = as.matrix(coef(cv.lasso, s = "lambda.min"))
```

Best: 23.606. This is better than both LS and Ridge. The corresponding lambda equals 0.025. Since this optimal lambda is small, only one variable has coefficient 0. This is INDUS, exactly the variable kicked out by model selection.

```
coefs <- data.frame(Lasso, Ridge, LM)
colnames(coefs) <- c("Lasso", "Ridge", "LM")
coefs
```

##	Lasso	Ridge	LM
## (Intercept)	37.44529	36.03344	39.40033
## CRIM	-0.10427	-0.10405	-0.11352
## ZN	0.05961	0.05756	0.06372
## INDUS	0.00000	-0.00258	0.02483
## CHAS	1.50770	1.62462	1.48609
## NOX	-15.67319	-15.02044	-17.03198
## RM	3.40370	3.48789	3.35805
## AGE	-0.00364	-0.00612	-0.00537
## DIS	-1.56345	-1.51390	-1.64248
## RAD	0.25312	0.23958	0.29926
## TAX	-0.01160	-0.01095	-0.01378
## PTRATIO	-0.82062	-0.81602	-0.84220

```
## B          0.00609  0.00641  0.00631
## LSTAT      -0.52102 -0.50413 -0.52157
```

9 All three methods

In the end, the difference in performance of the methods is immaterial.

```
set.seed(123)
loss <- function(predict){
  mean((predict- Test[,14])^2)
}
lm.model <- lm(MEDV ~., data = Train[, -which(names(Test) == "INDUS")])
predict.lm <- data.matrix(cbind(1, Test[, -c(which(names(Test) == "INDUS"),14)]))%*%
  data.matrix(lm.model$coef)
loss(predict.lm)
```

```
## [1] 25.2
```

```
x_test = as.matrix(Test[, c(1:13)])
predict.Ridge <- predict(cv.ridge, newx = x_test, s="lambda.min")
loss(as.matrix(predict.Ridge))
```

```
## [1] 25.1
```

```
predict.Lasso <- predict(cv.lasso, newx = x_test, s="lambda.min")
loss(as.matrix(predict.Lasso))
```

```
## [1] 25.2
```